

Joe Tsai

✉ joemhtsai@gmail.com 🌐 <https://cdes5804.github.io>

EDUCATION

Stanford University

M.S. in Computer Science

Sep. 2023 - Jun. 2025 (Expected)

Stanford, CA

- **GPA:** 4.22/4.30

National Taiwan University (NTU)

B.S. in Computer Science

Sep. 2018 - Jun. 2022

Taipei, Taiwan

- **GPA:** 4.28/4.30; **Rank:** 1/181; **Honors:** Phi Tau Phi (Honor for **top 1%** students)

WORK EXPERIENCE

Vectra AI

Software Engineer Intern

Jun. 2024 - Sep. 2024 (Expected)

San Jose, CA

- Migrated Kubernetes cron jobs to Celery Beat to eliminate the overhead associated with pod creation and teardown.
- Enhanced both UI and API responsiveness by transforming computation-intensive, synchronous threat detection processes into asynchronous background tasks, achieving near-instantaneous response times.

Academia Sinica

Research Assistant

Jun. 2022 - Dec. 2022

Taipei, Taiwan

- Implemented an intermittent system operating with scarce energy supply on STM32 devices in C, overcoming low-level programming challenges such as interrupt handling, race conditions, and memory I/O.
- Established Bluetooth Low Energy communication and optimized parameters to support data exchange between multiple devices.
- Developed algorithms for optimal intermittent computation and data exchange, achieving 30% less latency than existing methods.

Synology

Software Engineer Intern, Development Infrastructure Team

Jul. 2021 - Dec. 2021

New Taipei, Taiwan

- Developed a CI/CD pipeline, streamlining builds for 10+ iOS and Android apps while minimizing user's script-writing effort.
- Configured and set up LVM, Postfix, Samba, Docker images and containers for the build service on bare-metal servers.
- Rewrote legacy build scripts in PHP, Ruby, and Bash into reusable Python scripts to improve script maintainability.

Microsoft

Software Engineer Intern, Geocoding Team, AI Bing Department

Jul. 2020 - Jun. 2021

Taipei, Taiwan

- Increased the queries-per-second rate of a Bing Maps analysis tool by ~20% by identifying I/O bottleneck and using async I/O.
- Led the integration of a query parsing model and the analysis tool to replace heuristic parsing, allowing the tool to analyze location queries of different countries.
- Designed a pipeline to collect geographical information from terabytes of logs in C# and SQL, and perform data cleaning and processing with Python to create machine learning datasets for geocoding models.

SELECTED PROJECTS

Byzantine Fault-tolerant RAID-like Filesystem — C++, gRPC, FUSE

- Designed and Implemented a distributed, RAID-like filesystem capable of tolerating up to $\lfloor \frac{n-1}{2} \rfloor$ Byzantine faults.
- Implemented integration with FUSE, allowing existing applications using ordinary I/O operations to work with the filesystem.

Pintos — C

- Implemented core operating system functionalities, including thread scheduling, synchronization mechanisms, and virtual memory.
- Deepened understanding of the interaction between the kernel and user processes through the implementation of system calls.

NetCut — C++

- Developed an ARP spoofing tool capable of scanning users and disrupting the Internet connection of targeted devices.
- Utilized socket programming and multi-threading to manipulate the Internet connections of multiple targets concurrently.

TECHNICAL SKILLS

- **Programming Languages:** (Proficient) C, C++, Python; (Fluent) Go, JavaScript
- **Tools and frameworks:** Linux, Docker, Git, CI/CD (GitHub, GitLab), PySpark, Pytorch